

▼ Necessary preparations

```

1  import numpy as np
2  import pandas as pd
3  import torch
4
5  import PIL
6  print(PIL.PILLOW_VERSION)
7
8  train_on_gpu = torch.cuda.is_available()
9
10 if not train_on_gpu:
11     print('CUDA is not available.  Training on CPU ...')
12 else:
13     print('CUDA is available!  Training on GPU ...')
```



5.4.1
CUDA is available! Training on GPU ...

```

1  import pickle
2  import numpy as np
3  from skimage import io
4  import random
5
6  from tqdm import tqdm, tqdm_notebook
7  from PIL import Image
8  from pathlib import Path
9
10 from torchvision import transforms
11 from multiprocessing.pool import ThreadPool
12 from sklearn.preprocessing import LabelEncoder
13 from torch.utils.data import Dataset, DataLoader
14 import torch.nn as nn
15
16 from matplotlib import colors, pyplot as plt
17 %matplotlib inline
18
19 # в sklearn не все гладко, чтобы в colab удобно выводить картинки
20 # мы будем игнорировать warnings
21 import warnings
22 warnings.filterwarnings(action='ignore', category=DeprecationWarning)
```

```

1  SEED = 42
2
3  random.seed(SEED)
4  np.random.seed(SEED)
5  torch.manual_seed(SEED)
6  torch.cuda.manual_seed(SEED)
7  torch.backends.cudnn.deterministic = True
```

```
1  DATA MODES = ['train', 'val', 'test']
```

```


2 RESCALE_SIZE = 224
3 DEVICE = torch.device("cuda")

```

```

1 print(torch.__version__)

```

 1.3.0

▼ Dataset construction

```

1 class SimpsonsDataset(Dataset):
2     def __init__(self, files, mode, augmentations = None):
3         super().__init__()
4         self.files = files
5         self.mode = mode
6         self.augmentations = augmentations
7
8         if self.mode not in DATA_MODES:
9             print(f'wrong mode: {self.mode}')
10            raise NameError
11
12        self.len_ = len(self.files)
13        self.label_encoder = LabelEncoder()
14
15        if self.mode != 'test':
16            self.labels = [path.parent.name for path in self.files]
17            self.label_encoder.fit(self.labels)
18
19            with open('label_encoder.pkl', 'wb') as le_dump:
20                pickle.dump(self.label_encoder, le_dump)
21
22        def __len__(self):
23            return self.len_
24
25        def load_sample(self, file):
26            image = Image.open(file)
27            image.load()
28            return image
29
30        def __getitem__(self, index):
31            transform = transforms.Compose([
32                transforms.ToTensor(),
33                transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
34            ])
35            # трансформации для шума
36            custom_augmentations_01 = transforms.RandomOrder([
37                transforms.RandomHorizontalFlip(p=0.5),
38                transforms.RandomApply([transforms.RandomRotation(degrees=10)], p=0.25),
39                transforms.RandomApply([transforms.RandomResizedCrop(224, scale=(0.8, 1.2)], p=0.25),
40                transforms.RandomApply([transforms.RandomAffine((-10,10), (0.1,0.1))], p=0.25),
41                transforms.RandomPerspective(distortion_scale=0.1, p=0.25),
42                transforms.RandomApply([transforms.ColorJitter(brightness=0.02, contrast=0.02)], p=0.25)
43            ])

```

```

44
45     x = self.load_sample(self.files[index])
46     x = self._prepare_sample(x)
47     #x = np.array(x / 255, dtype='float32')
48
49
50     if self.mode == 'test':
51         x = np.array(x)
52         x = np.array(x / 255, dtype='float32')
53         x = transform(x)
54         return x
55     else:
56
57         if self.mode == 'train':
58             x = custom_augmentations_01(x)
59
60             x = transform(x)
61
62             label = self.labels[index]
63             label_id = self.label_encoder.transform([label])
64             y = label_id.item()
65             return x, y
66
67     def _prepare_sample(self, image):
68         image = image.resize((RESCALE_SIZE, RESCALE_SIZE))
69         return image #np.array(image)

```

```

1  TRAIN_DIR = Path('/kaggle/input/simpsons4/train')
2  TEST_DIR = Path('/kaggle/input/simpsons4/testset/testset')
3
4  train_val_files = sorted(list(TRAIN_DIR.rglob('*.jpg')))
5  test_files = sorted(list(TEST_DIR.rglob('*.jpg')))

```

Тренируем модель на всей выборке, чтобы точно задействовать все классы. При этом вали тех же данных, НО аугментируем мы только тренировочную часть.

```

1  from sklearn.utils import shuffle
2  from sklearn.model_selection import train_test_split
3
4  train_val_labels = [path.parent.name for path in train_val_files]
5  train_files, val_files = train_test_split(train_val_files, test_size=0.3, \
6                                             stratify=train_val_labels)
7
8  train_files = shuffle(train_val_files, random_state=0) #!!!
9
10 val_dataset = SimpsonsDataset(val_files, mode='val')
11 train_dataset = SimpsonsDataset(train_files, mode='train')

```

▼ Let's take a look at our data

```

1
2 def imshow(img, title=None, plt_ax=plt, default=False):
3     img = img.numpy().transpose((1, 2, 0))
4     mean = np.array([0.485, 0.456, 0.406])
5     std = np.array([0.229, 0.224, 0.225])
6     img = std * img + mean
7     img = np.clip(img, 0, 1)
8     plt_ax.imshow(img)
9     if title is not None:
10         plt_ax.set_title(title)
11         plt_ax.grid(False)
12
13 fig, ax = plt.subplots(nrows=3, ncols=3, figsize=(10,10), sharex=True, sharey=
14
15 for fig_x in ax.flatten():
16     random_characters = int(np.random.uniform(0,1000))
17     im_val, label = val_dataset[random_characters]
18     img_label = " ".join(map(lambda x: x.capitalize(),\
19                             val_dataset.label_encoder.inverse_transform([label])[0].split()
20     imshow(im_val.data.cpu(), \
21           title=img_label,plt_ax=fig_x)
22

```



▼ Building the model

Импортируем готовую resnet34 и заменим последний слой. Ничего не замораживаем.

```

1  from torchvision import models
2
3  model = models.resnet34(pretrained=True)
4  #model = models.resnet18(pretrained=True)

1  for param in model.parameters():
2      param.requires_grad = True
3
4  model.fc = nn.Linear(in_features=model.fc.in_features, out_features=42)

1  model = model.to(DEVICE)

```

Создадим циклический lr_scheduler, меняющий lr от 0.001 до 0.01:

```

1  import math
2
3  def cyclical_lr(stepsize, min_lr=1e-3, max_lr=1e-2):
4
5      # самая простая треугольная функция изменения lr
6      scaler = lambda x: 1.
7
8      # дополнительная функция, чтобы узнать, на каком этапе цикла мы находимся
9      def relative(it, stepsize):
10         cycle = math.floor(1 + it / (2 * stepsize))
11         x = abs(it / stepsize - 2 * cycle + 1)
12         return max(0, (1 - x)) * scaler(cycle)
13
14     # лямбда-функция, считающая текущий lr
15     lr_lambda = lambda it: min_lr + (max_lr - min_lr) * relative(it, stepsize)
16     print('cycle')
17     return lr_lambda

```

Функции для обучения, проверки, предсказаний.

```

1  def fit_epoch(model, train_loader, criterion, optimizer, scheduler):
2      running_loss = 0.0
3      running_corrects = 0
4      processed_data = 0
5
6      for inputs, labels in train_loader:
7          inputs = inputs.to(DEVICE)
8          labels = labels.to(DEVICE)
9          optimizer.zero_grad()
10

```

```

11         outputs = model(inputs)
12         loss = criterion(outputs, labels)
13         loss.backward()
14
15         optimizer.step()
16         scheduler.step() # потому что pytorch.__version__ > 1.1.0
17
18         preds = torch.argmax(outputs, 1)
19         running_loss += loss.item() * inputs.size(0)
20         running_corrects += torch.sum(preds == labels.data)
21         processed_data += inputs.size(0)
22
23     train_loss = running_loss / processed_data
24     train_acc = running_corrects.cpu().numpy() / processed_data
25     return train_loss, train_acc
26
27 def eval_epoch(model, val_loader, criterion):
28     model.eval()
29     running_loss = 0.0
30     running_corrects = 0
31     processed_size = 0
32
33     for inputs, labels in val_loader:
34         inputs = inputs.to(DEVICE)
35         labels = labels.to(DEVICE)
36
37         with torch.set_grad_enabled(False):
38             outputs = model(inputs)
39             loss = criterion(outputs, labels)
40             preds = torch.argmax(outputs, 1)
41
42             running_loss += loss.item() * inputs.size(0)
43             running_corrects += torch.sum(preds == labels.data)
44             processed_size += inputs.size(0)
45     val_loss = running_loss / processed_size
46     val_acc = running_corrects.double() / processed_size
47     return val_loss, val_acc
48
49 def train(train_files, val_files, model, epochs, batch_size):
50     train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
51     val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
52
53     history = []
54     log_template = "\nEpoch {ep:03d} train_loss: {t_loss:0.4f} \
55 val_loss {v_loss:0.4f} train_acc {t_acc:0.4f} val_acc {v_acc:0.4f}"
56
57     with tqdm(desc="epoch", total=epochs) as pbar_outer:
58
59         #*****
60         #opt = torch.optim.Adam(model.parameters())
61         #criterion = nn.CrossEntropyLoss()
62
63         #*****
64         criterion = nn.CrossEntropyLoss()
65         #opt = torch.optim.SGD(model.parameters(), lr=1.)

```

```

66     opt = torch.optim.Adam(model.parameters())
67     step_size = 10*len(train_loader)
68     clr = cyclical_lr(step_size)
69     clr_scheduler = torch.optim.lr_scheduler.LambdaLR(opt, [clr])
70     #*****
71
72     for epoch in range(epochs):
73         train_loss, train_acc = fit_epoch(model, train_loader, criterion,
74         print("loss", train_loss)
75
76         val_loss, val_acc = eval_epoch(model, val_loader, criterion)
77         history.append((train_loss, train_acc, val_loss, val_acc))
78
79         pbar_outer.update(1)
80         tqdm.write(log_template.format(ep=epoch+1, t_loss=train_loss,\
81                                     v_loss=val_loss, t_acc=train_acc, \
82
83     return history

```

```

1  def predict(model, test_loader):
2      with torch.no_grad():
3          logits = []
4
5          for inputs in test_loader:
6              inputs = inputs.to(DEVICE)
7              model.eval()
8              outputs = model(inputs).cpu()
9              logits.append(outputs)
10
11     probs = nn.functional.softmax(torch.cat(logits), dim=-1).numpy()
12     return probs

```

▼ Actual training

После 20 эпох лосс снова возрастает, поэтому остановимся на них.

```
1 history = train(train_dataset, val_dataset, model=model, epochs=20, batch_size=
```



```

epoch:  0%|          | 0/20 [00:00<?, ?it/s]cycle
loss 3.7535036235146544
epoch:  5%|          | 1/20 [03:58<1:15:40, 238.96s/it]
Epoch 001 train_loss: 3.7535      val_loss 3.3938 train_acc 0.0281 val_acc 0.09
loss 1.554902995133205
epoch: 10%|          | 2/20 [07:55<1:11:26, 238.13s/it]
Epoch 002 train_loss: 1.5549      val_loss 0.7445 train_acc 0.6279 val_acc 0.82
loss 0.6358150311669669
epoch: 15%|          | 3/20 [11:51<1:07:16, 237.45s/it]
Epoch 003 train_loss: 0.6358      val_loss 0.4181 train_acc 0.8419 val_acc 0.89
loss 0.4281362249820133
epoch: 20%|          | 4/20 [15:46<1:03:10, 236.88s/it]
Epoch 004 train_loss: 0.4281      val_loss 0.2862 train_acc 0.8909 val_acc 0.92

```

```
1 loss, acc, val loss, val acc = zip(*history)
```

Наконец, посмотрим на график:

```

1  loss, acc, val_loss, val_acc = zip(*history)
2  fig = plt.figure(figsize=(20, 15))
3
4  losses = fig.add_subplot(2,2,3)
5  losses.plot(loss, label="train_loss")
6  losses.plot(val_loss, label="val_loss")
7  losses.legend(loc='best')
8  losses.grid(axis = 'y')
9  losses.set_xlabel("epochs")
10 losses.set_ylabel("loss")
11
12 accs = fig.add_subplot(2,2,4)
13 accs.plot(acc, label="train_acc")
14 accs.plot(val_acc, label="val_acc")
15 accs.legend(loc='best')
16 accs.grid(axis = 'y')
17 accs.set_xlabel("epochs")
18 accs.set_ylabel("accuracy")
19
20 plt.show()
```

▼ Submission

```

1  label_encoder = pickle.load(open("label_encoder.pkl", 'rb'))
2
3  test_dataset = SimpsonsDataset(test_files, mode="test")
4  test_loader = DataLoader(test_dataset, shuffle=False, batch_size=64)
5  probs = predict(model, test_loader)
6
7  preds = label_encoder.inverse_transform(np.argmax(probs, axis=1))
8  test_filenames = [path.name for path in test_dataset.files]

1  submit = pd.DataFrame({'Id': test_filenames, 'Expected': preds})
2  submit

1  submit.to_csv('model-resnet34_train100_batch64_epoch20.csv', index=False)
```